

Group 10

Sudhay Senthilkumar, Sandeep Rajakrishnan

Amrita School of Engineering - Coimbatore

Botnet Threat Detection using Machine Learning Techniques

Sandeep Rajakrishnan, Sudhay Senthilkumar,
Senthilkumar Thangavel^{*1} and Sulakshan Vajipayajula^{*2}

Department of Computer Science and Engineering

^{*1}Amrita School of Engineering, Coimbatore,

Amrita Vishwa Vidyapeetham, India.

^{*2}Architect-CTO Office

IBM Security, Bangalore

sandur43@gmail.com ,
sudhay2001@gmail.com ,
t_senthilkumar@cb.amrita.edu ^{*1},
svajipay@in.ibm.com ^{*2}

```
pip install category_encoders
```

```
Requirement already satisfied: category_encoders in /usr/local/lib/python3.7/dist-packages (2.3.0)
Requirement already satisfied: statsmodels>=0.9.0 in /usr/local/lib/python3.7/dist-packages (from category_encoders) (0.10.2)
Requirement already satisfied: patsy>=0.5.1 in /usr/local/lib/python3.7/dist-packages (from category_encoders) (0.5.2)
Requirement already satisfied: pandas>=0.21.1 in /usr/local/lib/python3.7/dist-packages (from category_encoders) (1.1.5)
Requirement already satisfied: numpy>=1.14.0 in /usr/local/lib/python3.7/dist-packages (from category_encoders) (1.19.5)
Requirement already satisfied: scikit-learn>=0.20.0 in /usr/local/lib/python3.7/dist-packages (from category_encoders) (1.0.1)
Requirement already satisfied: scipy>=1.0.0 in /usr/local/lib/python3.7/dist-packages (from category_encoders) (1.4.1)
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/dist-packages (from pandas>=0.21.1->category_encoders) (2018.9)
Requirement already satisfied: python-dateutil>=2.7.3 in /usr/local/lib/python3.7/dist-packages (from pandas>=0.21.1->category_encoders) (2.7.3)
Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (from patsy>=0.5.1->category_encoders) (1.15.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.20.0->category_encoders) (2.1.0)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.20.0->category_encoders) (1.1.0)
```

```
import pandas as pd
import numpy as np
import category_encoders as ce
import matplotlib.pyplot as plt
import tensorflow as tf
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import StandardScaler
```

```
from google.colab import drive
drive.mount('/content/drive',force_remount=True)
```

```
Mounted at /content/drive
```

```
data=pd.read_csv(r"/content/drive/MyDrive/iot_botnet_dataset/features_having_most_influence_on_Botnet_IoT.csv")
```

```
# data=pd.read_csv(r"/content/drive/MyDrive/iot_botnet_dataset/features_having_most_influence_on_Botnet_IoT.csv")
data=pd.read_csv(r"/content/drive/MyDrive/Group 10 - IBM Project - Sandeep and Sudhay/Temp_Folders/Datasets/features_having_most_influence_on_Botnet_IoT.csv")
```

```
data.head()
```

	pkSeqID	proto	saddr	sport	daddr	dport	seq	stddev	N_IN_Conn_P_SrcIP	min	state_number	mean	N_IN_C
0	792371	udp	192.168.100.150	48516	192.168.100.3	80	175094	0.226784	100	4.100436	4	4.457383	
1	2056418	tcp	192.168.100.148	22267	192.168.100.3	80	143024	0.451998	100	3.439257	1	3.806172	

```
# encoder1 = ce.HashingEncoder(cols='saddr',n_components=5)
# data = encoder1.fit_transform(data)
# encoder2 = ce.HashingEncoder(cols='daddr',n_components=5)
# data = encoder2.fit_transform(data)
```

len(data)

733705

data.dtypes

pkSeqID	int64
proto	object
saddr	object
sport	object
daddr	object
dport	object
seq	int64
stddev	float64
N_IN_Conn_P_SrcIP	int64
min	float64
state_number	int64
mean	float64
N_IN_Conn_P_DstIP	int64
drate	float64
srate	float64
max	float64
attack	int64
category	object
subcategory	object
dtype:	object

```
df = data.drop(['pkSeqID','subcategory','attack','dport','sport'],axis=1)
```

```
#encoder1 = ce.HashingEncoder(cols='saddr',n_components=6)
```

```
encoder1 = ce.BinaryEncoder(cols=['saddr'],return_df=True)
df = encoder1.fit_transform(df)
```

```
encoder2 = ce.BinaryEncoder(cols=['daddr'],return_df=True)
df = encoder2.fit_transform(df)
```

```
proto_encoded = pd.get_dummies(data=df['proto'],drop_first=True)
```

```
df = pd.concat([df,proto_encoded],axis=1)
df.drop('proto',axis=1,inplace=True)
```

df.dtypes

saddr_0	int64
saddr_1	int64
saddr_2	int64
saddr_3	int64
saddr_4	int64
daddr_0	int64
daddr_1	int64
daddr_2	int64
daddr_3	int64
daddr_4	int64
daddr_5	int64
seq	int64
stddev	float64
N_IN_Conn_P_SrcIP	int64
min	float64
state_number	int64
mean	float64
N_IN_Conn_P_DstIP	int64
drate	float64
srate	float64
max	float64
category	object
icmp	uint8
ipv6-icmp	uint8
tcp	uint8
udp	uint8
dtype:	object

```
scaler = StandardScaler()
X = df.drop('category',axis=1)
y = df['category']
X = scaler.fit_transform(X)
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,train_size=0.8,random_state=109)
```

```
clf = LogisticRegression(random_state=0,multi_class='multinomial').fit(X_train, y_train)
```

/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:818: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,

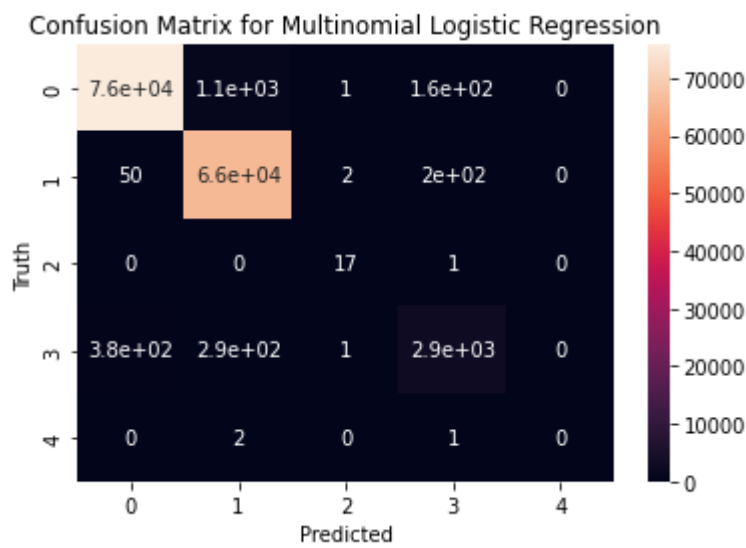
```
predictions = clf.predict(X_test)
print(accuracy_score(y_test,predictions))
```

0.9853347053652354

```
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
print("\t\tClassification report for Multinomial Logistic Regression\n\n",classification_report(y_test,predictions,digits=6))
print()
print()
plt.title("Confusion Matrix for Multinomial Logistic Regression")
sns.heatmap(confusion_matrix(y_test,predictions),annot=True)
plt.xlabel("Predicted")
plt.ylabel("Truth")
plt.show()
```

/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308: UndefinedMetricWarning: Precision and F-score are ill-def
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308: UndefinedMetricWarning: Precision and F-score are ill-def
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308: UndefinedMetricWarning: Precision and F-score are ill-def
_warn_prf(average, modifier, msg_start, len(result))
Classification report for Multinomial Logistic Regression

		precision	recall	f1-score	support
	DDoS	0.994369	0.984061	0.989188	76983
	DoS	0.979881	0.996145	0.987946	66152
	Normal	0.809524	0.944444	0.871795	18
	Reconnaissance	0.888584	0.814226	0.849782	3585
	Theft	0.000000	0.000000	0.000000	3
	accuracy			0.985335	146741
	macro avg	0.734472	0.747775	0.739742	146741
	weighted avg	0.985210	0.985335	0.985188	146741



```
import keras
from keras.models import Sequential
from keras.layers import Dense
```

```
from sklearn.preprocessing import OneHotEncoder
ohe = OneHotEncoder()
y = df['category'].values
```

```
y = ohe.fit_transform(y.reshape(-1,1)).toarray()
X_train,X_test,y_train,y_test = train_test_split(X,y,train_size=0.8,random_state=109)

model = Sequential()
model.add(Dense(16,input_dim=25,activation='relu'))
model.add(Dense(12,activation='relu'))
model.add(Dense(5,activation='softmax'))

model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

history = model.fit(X_train, y_train, epochs=100, batch_size=64)

Epoch 1/100
9172/9172 [=====] - 18s 2ms/step - loss: 0.0626 - accuracy: 0.9798
Epoch 2/100
9172/9172 [=====] - 19s 2ms/step - loss: 0.0254 - accuracy: 0.9901
Epoch 3/100
9172/9172 [=====] - 18s 2ms/step - loss: 0.0233 - accuracy: 0.9907
Epoch 4/100
9172/9172 [=====] - 17s 2ms/step - loss: 0.0219 - accuracy: 0.9914
Epoch 5/100
9172/9172 [=====] - 16s 2ms/step - loss: 0.0206 - accuracy: 0.9919
Epoch 6/100
9172/9172 [=====] - 17s 2ms/step - loss: 0.0194 - accuracy: 0.9925
Epoch 7/100
9172/9172 [=====] - 17s 2ms/step - loss: 0.0185 - accuracy: 0.9928
Epoch 8/100
9172/9172 [=====] - 18s 2ms/step - loss: 0.0175 - accuracy: 0.9932
Epoch 9/100
9172/9172 [=====] - 15s 2ms/step - loss: 0.0165 - accuracy: 0.9935
Epoch 10/100
9172/9172 [=====] - 16s 2ms/step - loss: 0.0157 - accuracy: 0.9940
Epoch 11/100
9172/9172 [=====] - 16s 2ms/step - loss: 0.0149 - accuracy: 0.9942
Epoch 12/100
9172/9172 [=====] - 17s 2ms/step - loss: 0.0143 - accuracy: 0.9945
Epoch 13/100
9172/9172 [=====] - 19s 2ms/step - loss: 0.0141 - accuracy: 0.9950
Epoch 14/100
9172/9172 [=====] - 19s 2ms/step - loss: 0.0127 - accuracy: 0.9951
Epoch 15/100
9172/9172 [=====] - 16s 2ms/step - loss: 0.0122 - accuracy: 0.9953
Epoch 16/100
9172/9172 [=====] - 17s 2ms/step - loss: 0.0119 - accuracy: 0.9954
Epoch 17/100
9172/9172 [=====] - 17s 2ms/step - loss: 0.0113 - accuracy: 0.9955
Epoch 18/100
9172/9172 [=====] - 17s 2ms/step - loss: 0.0116 - accuracy: 0.9957
Epoch 19/100
9172/9172 [=====] - 16s 2ms/step - loss: 0.0109 - accuracy: 0.9958
Epoch 20/100
9172/9172 [=====] - 16s 2ms/step - loss: 0.0105 - accuracy: 0.9959
Epoch 21/100
9172/9172 [=====] - 15s 2ms/step - loss: 0.0110 - accuracy: 0.9959
Epoch 22/100
9172/9172 [=====] - 16s 2ms/step - loss: 0.0103 - accuracy: 0.9959
Epoch 23/100
9172/9172 [=====] - 17s 2ms/step - loss: 0.0102 - accuracy: 0.9960
Epoch 24/100
9172/9172 [=====] - 17s 2ms/step - loss: 0.0098 - accuracy: 0.9961
Epoch 25/100
9172/9172 [=====] - 15s 2ms/step - loss: 0.0097 - accuracy: 0.9963
Epoch 26/100
9172/9172 [=====] - 16s 2ms/step - loss: 0.0094 - accuracy: 0.9962
Epoch 27/100
9172/9172 [=====] - 15s 2ms/step - loss: 0.0093 - accuracy: 0.9962
Epoch 28/100
9172/9172 [=====] - 16s 2ms/step - loss: 0.0114 - accuracy: 0.9963
Epoch 29/100
9172/9172 [=====] - 15s 2ms/step - loss: 0.0094 - accuracy: 0.9963
```

```
y_pred = model.predict(X_test)
```

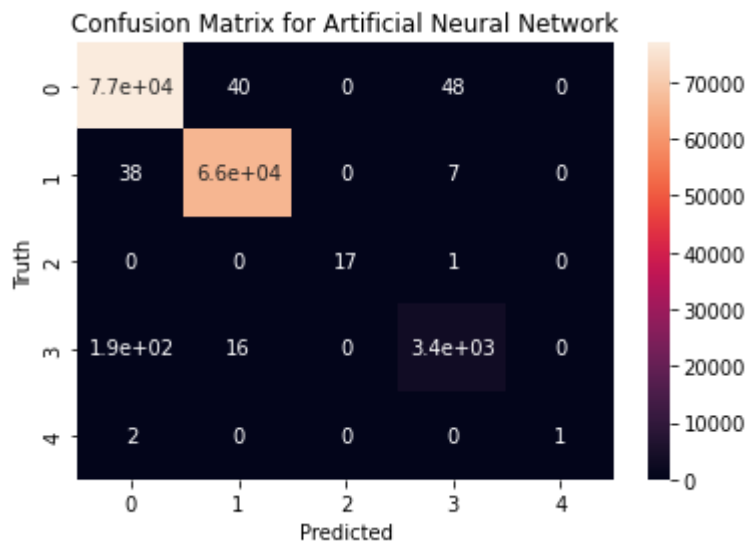
```
pred = list()
for i in range(len(y_pred)):
    pred.append(np.argmax(y_pred[i]))
test = list()
for i in range(len(y_test)):
    test.append(np.argmax(y_test[i]))
```

```
from sklearn.metrics import accuracy_score
a = accuracy_score(test,pred)
print('Accuracy is:', a)
```

Accuracy is: 0.9976966219393353

```
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
print("\t\tClassification report for Artificial Neural Network\n\n",classification_report(test,pred,digits=6))
print()
print()
plt.title("Confusion Matrix for Artificial Neural Network")
sns.heatmap(confusion_matrix(test,pred),annot=True)
plt.xlabel("Predicted")
plt.ylabel("Truth")
plt.show()
```

Classification report for Artificial Neural Network					
	precision	recall	f1-score	support	
0	0.997070	0.998857	0.997962	76983	
1	0.999154	0.999320	0.999237	66152	
2	1.000000	0.944444	0.971429	18	
3	0.983716	0.943654	0.963269	3585	
4	1.000000	0.333333	0.500000	3	
accuracy			0.997697	146741	
macro avg	0.995988	0.843922	0.886379	146741	
weighted avg	0.997683	0.997697	0.997676	146741	



Support Vector Machine - SVCClassifier

```
from sklearn.svm import SVC
model = SVC(kernel="linear")
model.fit(X_train,y_train)

print('Accuracy is:', model.score(X_test,y_test)*100)
```

GridSearch with SVC

```
from sklearn.model_selection import GridSearchCV
param_grid = {'C':(1, 10, 100, 1000), 'gamma':(0.1, 0.01, 0.001, 0.0001), 'kernel':('rbf','linear','poly')}
grid_search = GridSearchCV(model, param_grid, cv=10, verbose=10)
grid_search.fit(X_train, y_train)

grid_search.best_params_
```

Decision Tree Classifier with GridSearch

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import GridSearchCV
dec = DecisionTreeClassifier()
params = {'max_depth':(1,2,3,10,100,1000)}
grid_search_dec = GridSearchCV(dec, params, cv=5, verbose=10, n_jobs=-1)
grid_search_dec.fit(X_train, y_train)

Fitting 5 folds for each of 6 candidates, totalling 30 fits
GridSearchCV(cv=5, estimator=DecisionTreeClassifier(), n_jobs=-1,
```

param_grid={'max_depth': (1, 2, 3, 10, 100, 1000)}, verbose=10)


grid_search_dec.best_params_

{'max_depth': 1000}

y_pred = grid_search_dec.predict(X_test)

```
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
print("\t\tClassification report for DecisionTree\n\n",classification_report(y_test,y_pred,digits=6))
print()
print()
plt.title("Confusion Matrix for DecisionTree")
sns.heatmap(confusion_matrix(y_test,y_pred),annot=True)
plt.xlabel("Predicted")
plt.ylabel("Truth")
plt.show()
```

Classification report for DecisionTree					
		precision	recall	f1-score	support
	DDoS	0.999870	0.999935	0.999903	76983
	DoS	0.999819	0.999940	0.999879	66152
	Normal	1.000000	0.722222	0.838710	18
	Reconnaissance	1.000000	0.997768	0.998883	3585
	Theft	1.000000	1.000000	1.000000	3
	accuracy			0.999850	146741
	macro avg	0.999938	0.943973	0.967475	146741
	weighted avg	0.999850	0.999850	0.999847	146741



Random Forest Classifier

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV
randForest = RandomForestClassifier()
params = {'max_depth':(1,2,9,10),'n_estimators':(10,15,30)}
grid_search = GridSearchCV(randForest, params, cv=5, verbose=10)
grid_search.fit(X_train, y_train)
```

Fitting 5 folds for each of 12 candidates, totalling 60 fits
[CV 1/5; 1/12] START max_depth=1, n_estimators=10.....
[CV 1/5; 1/12] END max_depth=1, n_estimators=10;; score=0.854 total time= 3.8s
[CV 2/5; 1/12] START max_depth=1, n_estimators=10.....
[CV 2/5; 1/12] END max_depth=1, n_estimators=10;; score=0.796 total time= 3.9s
[CV 3/5; 1/12] START max_depth=1, n_estimators=10.....
[CV 3/5; 1/12] END max_depth=1, n_estimators=10;; score=0.832 total time= 4.3s
[CV 4/5; 1/12] START max_depth=1, n_estimators=10.....
[CV 4/5; 1/12] END max_depth=1, n_estimators=10;; score=0.751 total time= 8.4s
[CV 5/5; 1/12] START max_depth=1, n_estimators=10.....
[CV 5/5; 1/12] END max_depth=1, n_estimators=10;; score=0.852 total time= 7.3s
[CV 1/5; 2/12] START max_depth=1, n_estimators=15.....
[CV 1/5; 2/12] END max_depth=1, n_estimators=15;; score=0.857 total time= 8.4s
[CV 2/5; 2/12] START max_depth=1, n_estimators=15.....
[CV 2/5; 2/12] END max_depth=1, n_estimators=15;; score=0.852 total time= 8.2s
[CV 3/5; 2/12] START max_depth=1, n_estimators=15.....
[CV 3/5; 2/12] END max_depth=1, n_estimators=15;; score=0.861 total time= 7.4s
[CV 4/5; 2/12] START max_depth=1, n_estimators=15.....
[CV 4/5; 2/12] END max_depth=1, n_estimators=15;; score=0.840 total time= 6.7s
[CV 5/5; 2/12] START max_depth=1, n_estimators=15.....
[CV 5/5; 2/12] END max_depth=1, n_estimators=15;; score=0.817 total time= 5.1s
[CV 1/5; 3/12] START max_depth=1, n_estimators=30.....

```
[CV 1/5; 3/12] END max_depth=1, n_estimators=30; score=0.838 total time= 8.5s
[CV 2/5; 3/12] START max_depth=1, n_estimators=30; score=0.862 total time= 8.8s
[CV 2/5; 3/12] END max_depth=1, n_estimators=30; score=0.862 total time= 8.8s
[CV 3/5; 3/12] START max_depth=1, n_estimators=30; score=0.831 total time= 8.5s
[CV 3/5; 3/12] END max_depth=1, n_estimators=30; score=0.831 total time= 8.5s
[CV 4/5; 3/12] START max_depth=1, n_estimators=30; score=0.863 total time= 8.2s
[CV 4/5; 3/12] END max_depth=1, n_estimators=30; score=0.863 total time= 8.2s
[CV 5/5; 3/12] START max_depth=1, n_estimators=30; score=0.841 total time= 8.3s
[CV 5/5; 3/12] END max_depth=1, n_estimators=30; score=0.841 total time= 8.3s
[CV 1/5; 4/12] START max_depth=2, n_estimators=10; score=0.903 total time= 5.5s
[CV 1/5; 4/12] END max_depth=2, n_estimators=10; score=0.903 total time= 5.5s
[CV 2/5; 4/12] START max_depth=2, n_estimators=10; score=0.849 total time= 5.4s
[CV 2/5; 4/12] END max_depth=2, n_estimators=10; score=0.849 total time= 5.4s
[CV 3/5; 4/12] START max_depth=2, n_estimators=10; score=0.904 total time= 5.1s
[CV 3/5; 4/12] END max_depth=2, n_estimators=10; score=0.904 total time= 5.1s
[CV 4/5; 4/12] START max_depth=2, n_estimators=10; score=0.876 total time= 5.1s
[CV 4/5; 4/12] END max_depth=2, n_estimators=10; score=0.876 total time= 5.1s
[CV 5/5; 4/12] START max_depth=2, n_estimators=10; score=0.844 total time= 6.2s
[CV 5/5; 4/12] END max_depth=2, n_estimators=10; score=0.844 total time= 6.2s
[CV 1/5; 5/12] START max_depth=2, n_estimators=15; score=0.911 total time= 7.5s
[CV 1/5; 5/12] END max_depth=2, n_estimators=15; score=0.911 total time= 7.5s
[CV 2/5; 5/12] START max_depth=2, n_estimators=15; score=0.880 total time= 7.3s
[CV 2/5; 5/12] END max_depth=2, n_estimators=15; score=0.880 total time= 7.3s
[CV 3/5; 5/12] START max_depth=2, n_estimators=15; score=0.912 total time= 6.6s
[CV 3/5; 5/12] END max_depth=2, n_estimators=15; score=0.912 total time= 6.6s
[CV 4/5; 5/12] START max_depth=2, n_estimators=15; score=0.911 total time= 6.9s
[CV 4/5; 5/12] END max_depth=2, n_estimators=15; score=0.911 total time= 6.9s
[CV 5/5; 5/12] START max_depth=2, n_estimators=15; score=0.862 total time= 7.3s
[CV 5/5; 5/12] END max_depth=2, n_estimators=15; score=0.862 total time= 7.3s
[CV 1/5; 6/12] START max_depth=2, n_estimators=30; score=0.903 total time= 13.5s
[CV 1/5; 6/12] END max_depth=2, n_estimators=30; score=0.903 total time= 13.5s
[CV 2/5; 6/12] START max_depth=2, n_estimators=30; score=0.903 total time= 13.4s
[CV 2/5; 6/12] END max_depth=2, n_estimators=30; score=0.903 total time= 13.4s
[CV 3/5; 6/12] START max_depth=2, n_estimators=30; score=0.893 total time= 13.0s
[CV 3/5; 6/12] END max_depth=2, n_estimators=30; score=0.893 total time= 13.0s
[CV 4/5; 6/12] START max depth=2, n estimators=30; score=0.893 total time= 13.0s
```

```
y_pred = grid_search.predict(X_test)
```

```
grid_search.best_params_
```

```
{'max_depth': 10, 'n_estimators': 10}
```

```
from sklearn.metrics import classification_report, confusion_matrix
import seaborn as sns
print("\t\tClassification report for RandomForest\n\n",classification_report(y_test,y_pred,digits=6))
print()
print()
plt.title("Confusion Matrix for RandomForest")
sns.heatmap(confusion_matrix(y_test,y_pred),annot=True)
plt.xlabel("Predicted")
plt.ylabel("Truth")
plt.show()
```

Classification report for RandomForest					
		precision	recall	f1-score	support
	DDoS	0.999545	0.999662	0.999604	76983
	DoS	0.999562	0.999516	0.999539	66152
	Normal	1.000000	0.722222	0.838710	18
Reconnaissance		0.999163	0.998884	0.999024	3585
Theft		1.000000	1.000000	1.000000	3
accuracy				0.999543	146741
macro avg		0.999654	0.944057	0.967375	146741
weighted avg		0.999543	0.999543	0.999541	146741

